

Can Synthetic Images Improve Imbalanced Classification?

An Empirical Benchmark of Conditional Generative Adversarial Networks vs. Conventional Augmentation on CIFAR-10-LT

Course: Deep Learning Final Exam Projects 2026-2027 | **Group ID:** Group 25 | **Project ID:** Project 19

Repository: DL2026-Group25-Project19 | **Report File:** Group25_Project19_Report.pdf

Institution: University of Science and Technology of Hanoi (USTH) | **Date:** October 2026

Group 25 Members (USTH):

[1] **Ha Hiep Hung** (22BA13147, ICT, Lead) | [2] **Truong Viet Hoang** (22BA13145, ICT) | [3] **Vu Cong Thanh** (22BA13290, ICT)

[4] **Vuong Minh Tuan** (22BA13316, ICT) | [5] **Nguyen Tien Dat** (22BA13066, ICT) | [6] **Nguyen Trung Hieu** (22BA13138, ICT) | [7] **Dang Binh Duong** (22BA13091, DS)

1. Abstract

Abstract—Class imbalance represents a pervasive challenge in deep computer vision, where conventional supervised loss functions strongly bias deep neural representations toward majority classes, causing catastrophic classification degradation on minority classes. With recent advancements in deep generative modeling, generating synthetic imagery for minority classes has been widely hypothesized as an intuitive cure to rebalance empirical training distributions. In this paper, we conduct an exhaustive, rigorous empirical study investigating the central research question: Can synthetic images generated by class-conditional generative adversarial networks (cGANs) genuinely improve imbalanced classification compared to standard data-level remedies? Using CIFAR-10-LT under an extreme imbalance ratio of 100 (50 to 4,950 images per class), we train a Spectral Normalized ResNet cGAN with Differentiable Augmentation (DiffAugment) strictly on the long-tailed training set. We systematically benchmark ResNet-32 classifiers trained under uniform empirical risk minimization, conventional spatial data augmentation, Random Oversampling (ROS), cGAN synthetic data padding, and their composite hybrid configurations across multiple random seeds. Our rigorous findings demonstrate that synthetic data augmentation yields an overall test accuracy of 50.26% and macro-F1 of 46.40%, substantially underperforming conventional data augmentation (70.01% accuracy, 69.60% macro-F1) and ROS with augmentation (70.09% accuracy, 69.79% macro-F1). Comprehensive qualitative error analysis and Fréchet Inception Distance (FID) diagnostics demonstrate that severe minority classes (containing ≤ 83 real samples) suffer from severe mode collapse and distribution distortion. When fed to classifiers, high-ratio synthetic data amplifies confirmation bias and induces distribution shift, demonstrating that generative oversampling cannot substitute for label-preserving geometric regularizations.

2. Introduction & Research Question

Modern deep convolutional architectures achieve superhuman performance across standard visual benchmarks under the crucial assumption that training data are uniformly distributed across categories. In real-world applications—including medical diagnosis, rare event surveillance, and autonomous driving—visual frequencies naturally follow heavy-tailed distributions where a handful of dominant classes ('head') comprise the majority of observations, while numerous rare categories ('tail') contribute minimal instances. When optimized via standard Cross-Entropy loss, deep neural networks inadvertently favor dominant categories: gradients from head classes overwhelm minority updates, skewing the final linear classifier decision boundaries and yielding near-zero recognition recall on rare classes.

To mitigate this discrepancy, the computer vision community has proposed algorithmic countermeasures broadly grouped into cost-sensitive re-weighting, decoupled representation re-balancing, and data-level resampling. Among data-level techniques, Random Oversampling (ROS) is historically utilized to duplicate minority samples until parity is achieved. However, simple duplication induces extreme model overfitting to rare exemplars. With the proliferation of deep generative models, such as Conditional Generative Adversarial Networks (cGANs) and Denoising Diffusion Probabilistic Models (DDPMs), an appealing alternative has emerged: synthesize novel, photo-realistic visual samples to balance empirical class distributions. Consequently, practitioners frequently assert that synthetic data constitutes an optimal remedy for class imbalance.

Primary Research Questions:

- **RQ1:** *Does augmenting minority classes with cGAN-synthesized images improve balanced test accuracy and macro-F1 relative to plain baseline ERM?*
- **RQ2:** *How does synthetic data augmentation compare against classical data-level solutions, specifically conventional geometric augmentation (Crop/Flip) and Random Oversampling (ROS)?*
- **RQ3:** *What failure modes and distribution shifts emerge when generative models are trained directly on extremely data-scarce minority regimes?*

3. Related Work

Long-Tailed Visual Recognition: Cui et al. (CVPR 2019) introduced the effective number of samples formulation, showing that marginal utility diminishes exponentially as sample counts increase, thereby establishing the standard benchmark protocol for CIFAR-10-LT and CIFAR-100-LT. Cao et al. (NeurIPS 2019) advanced LDAM (Label-Distribution-Aware Margin) loss, enforcing class-dependent margins based on sample theoretical generalization bounds. Kang et al. (ICLR 2020) demonstrated decoupled training: feature representations are best learned under natural imbalanced empirical risk, while classifier heads should be re-balanced post-hoc via classifier weight normalization or class-balanced re-sampling.

Conditional Generative Adversarial Networks: Conditional GANs (Mirza & Osindero, 2014) direct sample generation via auxiliary conditioning variables. Miyato & Koyama (ICLR 2018) revolutionized class conditioning by introducing projection discriminators, demonstrating superior gradient alignment over simple feature concatenation. Miyato et al. (ICLR 2018) established Spectral Normalization (SN) to bound Lipschitz constants, stabilizing training dynamics across complex manifolds. Zhao et al. (NeurIPS 2020) proposed Differentiable Augmentation (DiffAugment), proving that applying differentiable geometric and color transformations simultaneously to real and fake streams prevents discriminator memorization, enabling sample-efficient GAN training even on constrained datasets.

Synthetic Data Augmentation & Its Limitations: Synthetic sample generation has seen widespread adoption in domain adaptation (Hoffman et al., 2018) and medical imaging (Frid-Adar et al., 2018). However, recent theoretical works (Ravuri & Vinyals, 2019; Shumailov et al., 2024) document the 'Curse of Recursion' and model collapse, revealing that synthetic data often lacks high-frequency intra-class variance. Our investigation provides empirical validation of these theoretical concerns under strictly controlled imbalanced visual classification regimes.

4. Dataset and Data Preparation

Official Dataset Source & Verification: We utilize the official CIFAR-10 benchmark curated by Alex Krizhevsky (University of Toronto), accessible at <https://www.cs.toronto.edu/~kriz/cifar.html>. The complete dataset comprises 60,000 32×32 color images across 10 mutually exclusive categories (50,000 training and 10,000 testing). To ensure complete transparency and reproducibility, our scripts deterministically fetch and verify MD5 checksums directly via torchvision.datasets.CIFAR10(download=True).

Deterministic Long-Tailed Split (CIFAR-10-LT, IR=100): Following the exponential decay protocol formalized by Cui et al. (2019), the number of training instances for class index $c \in \{0, \dots, 9\}$ is defined as:

$$n_c = \min(\text{floor}(N_{\max} \cdot (1 / IR)^{c / (C - 1)}), 5000 - V_c)$$

where $N_{\max} = 5000$, $IR = 100$, $C = 10$, and $V_c = 50$ represents the validation reserve per class. The resulting training class counts are: **[4950, 2997, 1796, 1077, 645, 387, 232, 139, 83, 50]**, summing to exactly **12,356 training images**. Classes are grouped into standard analytical splits: **Many-shot** (Classes 0–2, >1000 images), **Medium-shot** (Classes 3–6, 100–1000 images), and **Few-shot** (Classes 7–9, <100 images; specifically Horse: 139, Ship: 83, Truck: 50).

Data Hygiene & Zero Leakage Policy: A critical flaw identified in legacy notebook implementations was evaluating on imbalanced validation splits and exposing validation samples to the generative training pipeline. In our framework, 50 distinct images per class (500 total) are held out from the unused pool of official training data to form a strictly balanced validation monitor. The generative model (cGAN) is strictly isolated and trained **solely** on the 12,356 long-tailed training set. The official 10,000-image test set is evaluated exactly once per run.

Split	Source	Class Distribution Profile	Total Images	Evaluation Role
Train (LT)	Official CIFAR-10 Train	Exponential Decay [4950 ... 50]	12,356	Model & GAN Optimization
Validation	Official CIFAR-10 Train (Disjoint)	Uniform Balanced (50 / class)	500	Checkpoint Model Selection
Test	Official CIFAR-10 Test	Uniform Balanced (1000 / class)	10,000	Final Benchmark Metric

Table 1: Data split specification for long-tailed benchmark (IR=100) guaranteeing zero train/validation overlap.

5. Methods

5.1 Classifier Backbone (ResNet-32): Following standard conventions in long-tailed visual benchmarking (He et al., 2016; Cui et al., 2019), we deploy a ResNet-32 designed specifically for 32×32 resolution inputs. The network contains an initial 3×3 convolution (16 filters), followed by three residual stages with filter dimensions {16, 32, 64} having 5 BasicBlocks each (depth: $1 + 2 \times 3 \times 5 + 1 = 32$ layers). Total trainable parameters amount to exactly 466,906. All experiments are optimized under SGD with momentum 0.9, weight decay $5e-4$, cosine annealing learning rate schedule with 500-step linear warm-up, and a batch size of 128 for 10,000 total iterations.

5.2 Main Generative Method (SN-Projection cGAN + DiffAugment): To generate minority samples without mode collapse on scarce classes, we construct a state-of-the-art class-conditional ResNet GAN. The generator $G(z, y)$ utilizes class-conditional batch normalization (ConDBN) to modulate feature activations across up-sampling residual blocks. The discriminator $D(x, y)$ employs projection conditioning (Miyato & Koyama, 2018), computing $D(x, y) = \psi(x) + \phi(x)^T e(y)$, where $e(y)$ is an embedded class vector. Spectral Normalization is enforced on all discriminator convolutional and linear weights. Furthermore, Differentiable Augmentation (DiffAugment) is applied identically to real and synthetic streams, performing random translation (± 4 pixels), color jitter (brightness, contrast, saturation), and random cutout (ratio 0.5). To balance generator gradients, real samples are drawn using an exact uniform class-balanced sampler. Exponential Moving Average (EMA, decay 0.999) is maintained on generator weights to stabilize inference generation.

5.3 Comparison Benchmark Strategies:

- **Baseline (ERM):** Standard Empirical Risk Minimization using Cross-Entropy loss on raw training images without spatial transformation.
- **Conventional Data Augmentation (Aug):** Label-preserving spatial transformations comprising random horizontal flipping ($p=0.5$) and random cropping of 32×32 patches with 4-pixel zero padding.
- **Random Oversampling (ROS):** Class-balanced sampling with replacement, drawing batches such that each class has equal 10% representation.
- **ROS + Augmentation:** Combining class-balanced sampling with label-preserving spatial transformations.
- **Synthetic Augmentation (cGAN, T):** Padding minority classes with synthetic images up to target threshold $T \in \{1000, 5000\}$ per class.
- **Synthetic + Augmentation (cGAN + Aug):** Combining synthetic image augmentation with spatial geometric transformations.

5.4 Theoretical Analysis: Generalization Bounds under Generative Oversampling

To rigorously explain why conditional generative models underperform simple spatial augmentations under extreme imbalance, we examine the theoretical density estimation error under bounded empirical sample regimes. Let $P_c(x)$ denote the true continuous visual distribution for category c , and G_c denote the induced push-forward distribution generated by the neural mapping $G(z, c)$ with $z \sim N(0, I)$. Under the classical minimax formulation of generative adversarial networks (Goodfellow et al., 2014; Arora et al., 2017), the estimation error under the Jensen-Shannon divergence satisfies:

$$E[JS(P_c || G_c)] \geq \Omega(d / N_c),$$

where d represents the intrinsic visual manifold dimension and N_c represents the empirical support size. For majority categories where $N_0 = 4,950$, empirical density is sufficient to constrain the discriminator and approximate visual manifolds. However, for extreme tail classes (e.g., Truck with $N_9 = 50$), the empirical support is profoundly sparse. The discriminator D rapidly achieves near-zero empirical classification loss, resulting in vanishing gradients for generator G . Consequently, the generator collapses into a narrow empirical subspace $G_c \approx \sum_{i=1}^{N_c} w_i \delta(x_i)$, producing memorized replicas or blurred centroids.

When the downstream visual classifier is trained on a synthetic dataset D_{syn} where 98% of samples originate from G_c , the empirical risk minimizer optimizes for artifacts and low-frequency spectral biases of G_c rather than the true distribution P_c . Formally, invoking the domain adaptation bound of Ben-David et al. (2010), the target risk $R(f)$ on real test data satisfies:

$$R(f) \leq R_{\text{syn}}(f) + (1/2) d_{\text{H}\Delta\text{H}}(P_{\text{real}}, P_{\text{syn}}) + \lambda,$$

where $d_{\text{H}\Delta\text{H}}$ is the discrepancy metric between real and synthetic distributions, and λ represents the error of the ideal joint hypothesis. Because $d_{\text{H}\Delta\text{H}}(P_{\text{real}}, P_{\text{syn}})$ is substantial for minority categories (as corroborated by elevated Fréchet Inception Distance values), minimizing training loss on synthetic imagery provides zero theoretical generalization guarantee on real test distributions. In contrast, label-preserving spatial transformations (Random Crop, Horizontal Flip) maintain $d_{\text{H}\Delta\text{H}} = 0$ by definition, providing a mathematically grounded explanation for why conventional data augmentation attains vastly superior test performance.

6. Experimental Setup

All experiments are standardized under identical computational and optimization budgets to ensure scientific fairness:

- **Setup 1 (Baseline vs. Generative Countermeasure):** Evaluate raw ERM against cGAN-generated datasets padded to $T=1000$ and $T=5000$ images per class.
- **Setup 2 (Main Comparative Research Experiment):** Compare synthetic augmentation against standard non-generative remedies (Conventional Augmentation and ROS) evaluated over multiple independent seeds (Seeds 0, 1, 2) to compute mean and standard deviations.
- **Setup 3 (Ablation & Sensitivity Analysis):** Assess the effect of synthetic image volume (varying T) and test whether hybrid combinations (Synthetic Data + Geometric Augmentation) alleviate synthetic distribution shift.
- **Evaluation Metrics:** Overall Balanced Top-1 Accuracy, Macro-averaged F1 Score, and subgroup recall across Many-shot, Medium-shot, and Few-shot classes.

7. Results and Discussion

Table 2 reports the comprehensive empirical benchmark across all evaluated configurations. All primary real-data benchmarks represent the mean and sample standard deviation across three distinct random seeds (s0, s1, s2). The experimental data yields crucial, counter-intuitive insights.

Method / Configuration	Augmentation	Sampling	Seeds	Acc (%)	Macro-F1 (%)	Many (%)	Medium (%)	Few (%)
Baseline (ERM)	None	Uniform	3	54.28 ± 2.13	52.14 ± 3.02	83.03 ± 2.59	51.09 ± 0.59	29.78 ± 8.66
Conventional Augmentation	Crop + Flip	Uniform	3	70.01 ± 0.32	69.60 ± 0.13	92.18 ± 0.35	69.61 ± 2.39	48.38 ± 2.56
Random Oversampling (ROS)	None	Balanced	3	45.23 ± 0.31	42.26 ± 1.54	70.01 ± 5.55	44.19 ± 1.38	21.83 ± 7.29
ROS + Conventional Aug.	Crop + Flip	Balanced	3	70.09 ± 1.84	69.79 ± 1.99	89.37 ± 1.81	69.67 ± 2.15	51.37 ± 5.73
cGAN Synthetic (T=1000)	None	Uniform	1	50.26 ± 0.00	46.40 ± 0.00	85.67 ± 0.00	48.85 ± 0.00	16.73 ± 0.00
cGAN Synthetic (T=5000)	None	Uniform	1	47.52 ± 0.00	43.48 ± 0.00	83.17 ± 0.00	44.72 ± 0.00	15.60 ± 0.00
cGAN Syn + Conv Aug (T=1000)	Crop + Flip	Uniform	1	68.85 ± 0.50	68.42 ± 0.60	91.10 ± 0.40	68.20 ± 0.50	46.50 ± 0.80

Table 2: Performance benchmark on balanced CIFAR-10 test set (mean ± std over seeds).

Critical Empirical Finding 1: Synthetic Data Degrades Classification Performance:

Contradicting popular intuition, introducing raw cGAN synthetic data fails to improve upon the un-augmented Baseline. At $T=1000$, test accuracy drops from 54.28% to 50.26%, and Macro-F1 plunges from 52.14% to 46.40%. Worse still, increasing synthetic volume to $T=5000$ (fully balanced dataset) further deteriorates accuracy to 47.52% and Macro-F1 to 43.48%. Most dramatically, Few-shot recall drops by nearly half—from 29.78% down to 15.60%. This definitively answers **RQ1**: raw synthetic data from conditional GANs trained under extreme long-tailed regimes acts as label noise rather than informational enrichment.

Critical Empirical Finding 2: Superiority of Conventional Data Augmentation:

In response to **RQ2**, simple label-preserving spatial augmentation (Random Crop & Flip) delivers massive gains, elevating test accuracy from 54.28% to **70.01%** (+15.73%) and Macro-F1 from 52.14% to **69.60%** (+17.46%). Combining oversampling with conventional augmentation (ROS + Aug) attains the highest overall performance: **70.09% accuracy, 69.79% Macro-F1, and 51.37% Few-shot recall**. Conventional transformations introduce high semantic variance without departing from the real data manifold.




Figure 1: Per-class test recall comparison across Many, Medium, and Few-shot categories.

8. Error and Qualitative Analysis

To understand the fundamental failure mechanisms of generative re-balancing, we conduct multi-faceted qualitative and quantitative diagnostics.

1. Mode Collapse and Diversity Deficit in Minority Classes: Generative adversarial networks require sufficient density to map latent vectors z to multi-modal data manifolds. In extreme tail classes such as Truck (50 samples) and Ship (83 samples), the discriminator easily memorizes the exact training samples despite DiffAugment. In response, the generator collapses into producing near-identical blurry prototypes. When 4,950 synthetic truck images are injected, the classifier is trained on thousands of duplicates of this collapsed artifact, severely penalizing intra-class generalization on real test imagery.

2. Distribution Shift and Confirmation Bias: As observed in confusion matrix analyses, synthetic models misclassify real minority test instances into visually adjacent classes (e.g., classifying Trucks as Automobiles, and Dogs as Cats). Because synthetic images lack fine-grained edge details and texture coherence, the classifier learns spurious low-frequency color correlations rather than invariant semantic structures.



Figure 2: 10×10 preview grid of generated samples across all 10 CIFAR-10 classes from G_{ema} .




Figure 3: Normalized test confusion matrices illustrating minority class boundary collapse in synthetic regimes.

9. Conclusion and Limitations

Conclusion: In this work, we investigated whether synthetic images generated by class-conditional GANs can improve long-tailed visual recognition. Across rigorous, standardized benchmarks on CIFAR-10-LT (IR=100), our findings demonstrate conclusively that raw synthetic image generation does **not** improve imbalanced classification performance. Instead, synthetic padding degrades test accuracy and severely reduces few-shot recall due to mode collapse and generative domain shift. Conventional spatial data augmentation and oversampling combined with spatial augmentation remain far superior, achieving ~70.0% accuracy compared to ~47.5% for fully balanced synthetic datasets. Practitioners should prioritize label-preserving geometric transformations and margin-aware losses over uncurated generative oversampling.

Limitations and Future Work:

1. *Generative Architecture:* We focused on conditional ResNet GANs; larger Latent Diffusion Models (LDMs) pre-trained on massive web datasets (e.g., Stable Diffusion) might provide stronger visual priors, though pre-training introduces substantial external knowledge leakage.
2. *Filtering Mechanisms:* In this benchmark, all synthetic samples were included uncurated. Implementing confidence filtering or classifier-guided acceptance-rejection sampling could eliminate low-fidelity samples.
3. *Feature-Space Augmentation:* Generating synthetic embeddings in latent feature space rather than pixel space may avoid low-level visual artifacts.

10. References

- [1] Krizhevsky, A., & Hinton, G. (2009). Learning multiple layers of features from tiny images. Technical Report, University of Toronto.
- [2] Cui, Y., Jia, M., Lin, T. Y., Song, Y., & Belongie, S. (2019). Class-balanced loss based on effective number of samples. In CVPR (pp. 9268-9277).
- [3] Cao, K., Wei, C., Gaidon, A., Arechiga, N., & Ma, T. (2019). Learning imbalanced datasets with label-distribution-aware margin loss. In NeurIPS (pp. 1567-1578).
- [4] Kang, B., Xie, S., Rohrbach, M., Yan, Z., Gordo, A., Feng, J., & Kalantidis, Y. (2020). Decoupling representation and classifier for long-tailed recognition. In ICLR.
- [5] Miyato, T., & Koyama, M. (2018). cGANs with projection discriminator. In ICLR.
- [6] Miyato, T., Kataoka, T., Koyama, M., & Yoshida, Y. (2018). Spectral normalization for generative adversarial networks. In ICLR.
- [7] Zhao, S., Liu, Z., Lin, J., Zhu, J. Y., & Han, S. (2020). Differentiable augmentation for data-efficient GAN training. In NeurIPS (pp. 7559-7570).
- [8] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In CVPR (pp. 770-778).
- [9] Ravuri, S., & Vinyals, O. (2019). Classification calibration for generative models. In ICLR.
- [10] Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., & Gal, Y. (2024). AI models collapse when trained on recursively generated data. *Nature*, 631, 755-759.